

Sprint 2 - Steps & Approach

#	AI FDE Sprint Step	What Is Actually Delivered
1	Frame the Decision, Risk & Regulatory Boundary	Define the business decision, users affected, APPROVE/REVIEW/REJECT outcomes, anomaly/fraud scenarios, risk appetite, applicable regulatory obligations, prohibited automation and mandatory human-oversight points.
2	Engineer the Case, Data & Evidence Model	Define case events, entities, identity attributes, risk variables, labels, document fields, data-quality rules, lineage and the evidence that must be retained for every decision.
3	Build & Independently Evaluate Detection Models	Train anomaly/fraud models using representative labelled data; evaluate precision/recall, AUC, catch rate, false positives and segment performance; compare model versions and document limitations before deployment.
4	Extract Documents + Verify Identity & Compliance Signals	Convert incoming PDFs/forms/images into structured case data; validate identity consistency and confidence; integrate sanctions/PEP/address or other verification controls where applicable; route uncertain evidence for review.
5	Orchestrate the End-to-End Case Workflow	Connect Ingest → Extract → Verify → Detect → Score → Decide → Review → Notify, including parallel checks, state persistence, retries, timeouts, fallback paths and compensating actions.
6	Combine AI + Rules into Governed Risk Decisioning	Fuse anomaly scores, identity evidence, business rules, policy constraints and contextual risk into

#	AI FDE Sprint Step	What Is Actually Delivered
		APPROVE / REVIEW / REJECT, with policy thresholds and human-in-the-loop escalation rather than allowing the model to make uncontrolled final decisions.
7	Make Decisions Explainable, Reviewable & Contestable	Produce machine-readable reasons, human-readable rationale, feature attribution/counterfactual explanations where suitable, reviewer evidence, customer notification and mechanisms for human reconsideration or appeal where required.
8	Embed Compliance Monitoring + Immutable Audit Evidence	Detect velocity, behavioural, geographic and other relevant exceptions; generate compliance escalation events where applicable; record model version, inputs, rules, prompts, workflow states, human interventions and decisions in tamper-resistant/WORM evidence stores.
9	Operationalize AI Risk, Security & Observability	Establish model ownership, model cards, validation evidence, change approvals, access controls, drift/performance monitoring, false-positive trends, data-quality monitoring, audit completeness, latency, failure rates and compliance-control alerts.
10	Prove, Govern & Productionize the Complete AI Product	Execute functional, integration, adverse, failure, compliance and performance tests; validate fallback/human-review paths; resolve defects; package governance evidence; then establish CI/CD, IaC, release gates, rollback, security hardening and production handover.